

011000100110100101110100011100110111
bitstring
010001110010011010010110111001100111

A Python module to help you manage your bits

This document describes version 0.4.3 of the bitstring module.
To download the latest version go to <http://python-bitstring.googlecode.com>.

9 June 2009

The MIT License

Copyright © 2009 Scott Griffiths (scott@griffiths.name)

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

Contents

Introduction	3
Getting Started	3
Creation and Interpretation	4
Using the constructor	4
Interpreting BitStrings	7
Slicing, Dicing and Splicing	9
Slicing	9
Joining	10
Truncating	10
Inserting, deleting and overwriting	10
Splitting	12
A BitString is a list	11
Reading and Navigating	12
Reading	12
Seeking	13
Finding	13
Miscellany	14
Other Functions	14
Special Methods	15
Reference	18
BitString methods and properties	18
BitString special methods	23
Module functions	25
Exponential-Golomb Codes	26
Internals	28

1. Introduction

bitstring is a pure Python module designed to help make the creation and analysis of binary data as painless as possible.

It provides a single class, `BitString`, instances of which can be constructed from integers, hex, octal, binary, strings or files, but they all just represent a string of binary digits. They can be sliced, joined, reversed, inserted into, overwritten, etc. with simple functions or slice notation. They can also be read from, searched in, and navigated in, similar to a file or stream.

`BitString` objects are designed to be as lightweight as possible and can be considered to be just a list of binary digits. They are however stored very efficiently - although there are a variety of ways of creating and viewing the binary data, the `BitString` itself just stores the byte data, and all views are calculated as needed, and are not stored as part of the object.

The different views or interpretations on the data are accessed through properties such as `hex`, `bin` and `int`, and an extensive set of functions is supplied for modifying, navigating and analysing the binary data.

Just a flavour is given here, and will be covered in greater detail in the rest of this manual.

```
from bitstring import BitString

# Just some of the ways to create BitStrings
a = BitString('0b001')           # from a binary string
b = BitString('0xff470001')       # from a hexadecimal string
c = BitString(filename='somefile.ext') # straight from a file
d = BitString(int=540, length=11) # from an integer

# Easily construct new BitStrings
e = 5*a + '0xcdcd'               # 5 copies of 'a' followed by two new
                                # bytes
e.prepend('0b1')                 # put a single bit on the front
f = e[7:]                        # cut the first 7 bits off
f[1:4] = '0o775'                 # replace 3 bits with 9 bits from
                                # octal string
f.replace('0b01', '0xee34', False) # find and replace 2 bit string with
                                # 16 bit string

# Interpret the BitString however you want
print e.hex                      # 0x9249cdcd
print e.int                      # -1840656947 (signed 2's complement
                                # integer)
print e.uint                    # 2454310349 (unsigned integer)
open('somefile.ext', 'rb').write(e.data) # Output raw byte data to a file
```

A complete reference for the module is given in Appendix A.

1.1. Getting Started

First download the latest release for either Python 2.4 / 2.5 / 2.6 or 3.0 (see the Downloads tab on the project's [homepage](#)) and extract the contents of the zip. You should find:

- ▶ `bitstring.py` : The bitstring module itself.
- ▶ `test_bitstring.py` : Unit tests for the module.
- ▶ `setup.py` : The setup script.
- ▶ `readme.txt` : A short readme.
- ▶ `release_notes.txt` : History of changes in this and previous versions.

- ▶ `test/test.m1v` : An example file (MPEG-1 video) for testing purposes.
- ▶ `test/smalltestfile` : Another small file for testing.
- ▶ `docs/bitstring.pdf` : This document

To install, run

```
python setup.py install
```

This will copy `bitstring.py` to your Python installation's `site-packages` directory. If you prefer you can do this by hand, or just make sure that your Python program can see `bitstring.py`, for example by putting in the same directory as the program that will use it.

The module comes with comprehensive unit tests. To run them yourself use

```
python test_bitstring.py
```

which should run all the tests (over 200) and say **OK**. If tests fail then either your version of Python isn't supported (there's one version of `bitstring` for Python 2.4, 2.5 and 2.6 and a separate version for Python 3.0) or something unexpected has happened - in which case please tell me about it.

2. Creation and Interpretation

You can create `BitString` objects in a variety of ways. Internally, `BitString` objects are stored as byte arrays. This means that no space is wasted and a `BitString` containing 10MB of binary data will only take up 10MB of memory. When a `BitString` is created all that is stored is the byte array, the length in bits and possibly an offset to the first used bit in the byte array. This means that the actual initialiser used to create the `BitString` isn't stored itself - if you create using a hex string for example then if you ask for the hex interpretation it has to be calculated from the stored byte array.

2.1. Using the constructor

When initialising a `BitString` you need to specify at most one initialiser. These will be explained in full below, but briefly they are:

- `auto` : Either a string prefixed with `0x`, `0o` or `0b` to interpret as hexadecimal, octal or binary, another `BitString` or a list or tuple.
- `data` : A Python string, for example read from a binary file.
- `hex`, `oct`, `bin`: Hexadecimal, octal or binary strings.
- `int`, `uint`: Signed or unsigned binary integers.
- `se`, `ue` : Signed or unsigned exponential-Golomb coded integers.
- `filename` : Directly from a file, without reading into memory.

The auto initialiser

```
>>> fromhex = BitString('0x01ffc9')
>>> frombin = BitString('0b01')
>>> fromoct = BitString('0o7550')
>>> acopy = BitString(fromoct)
>>> fromlist = BitString([True, False, False])
```

The simplest way to create a `BitString` is often to use the `auto` parameter, which is the first parameter in the `__init__` function and so the `auto=` can be omitted. It accepts a number of different objects. Strings that start with `0x` are interpreted as hexadecimal, `0o` implies octal, and strings starting with `0b` are

interpreted as binary. If given another `BitString` it will create a copy of it. Lists and tuples are interpreted as boolean arrays.

Note that as always the `BitString` doesn't know how it was created. Initialising with octal or hex might be more convenient or natural for a particular example but it is exactly equivalent to initialising with the corresponding binary string.

```
>>> fromoct.oct
'0o7550'
>>> fromoct.hex
'0xf68'
>>> fromoct.bin
'0b111101101000'
>>> fromoct.uint
3994
>>> fromoct.int
-152
>>> BitString('0o7777') == '0xffff'
True
>>> BitString('0xf') == '0b1111'
True
>>> frombin[::-1] + '0b0' == fromlist
True
```

Note how in the final examples above only one half of the `==` needs to be a `BitString`, the other half gets 'auto' initialised. This is in common with many other functions and operators.

From raw data

For most initialisers you can also use the `length` and `offset` parameters to specify the length in bits and an offset at the start to be ignored. This is particularly useful when initialising from raw data for from a file.

```
a = BitString(data='\x00\x01\x02\xff', length=28, offset=1)
b = BitString(data=open("somefile", 'rb').read())
```

The `length` parameter is optional; it defaults to the length of the data in bits (and so will be a multiple of 8). You can use it to truncate some bits from the end of the `BitString`. The `offset` parameter is used to truncate bits at the start of the data.

From a file

Using the `filename` initialiser allows a file to be analysed without the need to read it all into memory. The way to create a file-based `BitString` is:

```
p = BitString(filename="my2GBfile")
```

This will open the file in binary read-only mode. The file will only be read as and when other operations require it, and the contents of the file will not be changed by any operations. If only a portion of the file is needed then the `offset` and `length` parameters (specified in bits) can be used.

Something to watch out for are operations that could cause a copy of large parts of the object to be made in memory, for example

```
p2 = p[8:]
p += '0x00'
```

will create two new memory-based `BitString` objects with about the same size as the whole of the file's data. This is probably not what is wanted as the reason for using the `filename` initialiser is likely to be because you don't want the whole file in memory.

From a hexadecimal string

```
c = BitString(hex='0x000001b3')
```

The initial `0x` or `0X` is optional, as once again is a `length` parameter, which can be used to truncate the end. Whitespace is also allowed and is ignored. Note that the leading zeros are significant, so the length of `c` will be 32.

If you include the initial `0x` then you can use the auto initialiser instead. As it is the first parameter in `__init__` this will work equally well:

```
c = BitString('0x000001b3')
```

From a binary string

```
>>> d = BitString(bin='0011 000', length=6)
>>> d.bin
'0b001100'
```

An initial `0b` or `0B` is optional. Once again a `length` can optionally be supplied to truncate the `BitString` and whitespace will be ignored.

As with hex, the auto initialiser will work if the binary string is prefixed by `0b`:

```
>>> d = BitString('0b001100')
```

From an octal string

```
>>> o = BitString(oct='34100')
>>> o.oct
'0o34100'
```

An initial `0o` or `0O` is optional, but `0o` (a zero and lower-case 'o') is preferred as it is slightly more readable. Once again a `length` can optionally be supplied to truncate the `BitString` and whitespace will be ignored.

As with hex and bin, the auto initialiser will work if the octal string is prefixed by `0o`:

```
>>> o = BitString('0o34100')
```

From an integer

```
>>> e = BitString(uint=45, length=12)
>>> f = BitString(int=-1, length=7)
>>> e.bin
'0b000000101101'
>>> f.bin
'0b1111111'
```

For initialisation with signed and unsigned binary integers (`int` and `uint` respectively) the `length` parameter is mandatory, and must be large enough to contain the integer. So for example if `length` is 8 then `uint` can be in the range 0 to 255, while `int` can range from -128 to 127. Two's complement is used to represent negative numbers.

Initialisation with integers represented by exponential-Golomb codes is also possible. `ue` is an unsigned code while `se` is a signed code:

```
>>> g = BitString(ue=12)
>>> h = BitString(se=-402)
>>> g.bin
'0b0001101'
>>> h.bin
'0b0000000001100100101'
```

For these initialisers the length of the `BitString` is fixed by the value it is initialised with, so the `length` parameter must not be supplied and it is an error to do so. If you don't know what exponential-Golomb codes are then you are in good company, but they are quite interesting, so I've included an appendix on them (see Appendix B).

2.2. Interpreting BitStrings

`BitString` objects don't know or care how they were created; they are just collections of bits. This means that you are quite free to interpret them in any way that makes sense.

Several Python properties are used to create interpretations for the `BitString`. These properties call private functions which will calculate and return the appropriate interpretation. These don't change the `BitString` in any way and it remains just a collection of bits. If you use the property again then the calculation will be repeated.

For the properties described below we will use these:

```
>>> a = BitString('0x123')
>>> b = BitString('0b111')
```

bin

The most fundamental interpretation is perhaps as a binary string (a 'bitstring'). The `bin` property returns a string of the binary representation of the `BitString` prefixed with `0b`. All `BitString` objects can use this property and it is used to test equality between `BitString` objects.

```
>>> a.bin
'0b000100100011'
>>> b.bin
'0b111'
```

Note that the initial zeros are significant; for `BitString` objects the zeros just as important as the ones!

hex

For whole-byte `BitString` objects the most natural interpretation is often as hexadecimal, with each byte represented by two hex digits. Hex values are prefixed with `0x`.

If the `BitString` does not have a length that is a multiple of four bits then a `ValueError` exception will be raised. This is done in preference to truncating or padding the value, which could hide errors in user code.

The `hex` built-in function can also be used, with exactly the same effect.

```
>>> a.hex
'0x123'
>>> hex(a)
'0x123'
>>> b.hex
ValueError: Cannot convert to hex unambiguously - not multiple of 4 bits.
```

oct

For an octal interpretation use the `oct` property or the `oct` built-in function. Octal values are prefixed with `0o`, which is the Python 2.6 / 3.0 way of doing things (rather than just starting with `0`).

If the `BitString` does not have a length that is a multiple of three then a `ValueError` exception will be raised.

```
>>> a.oct
'0o0443'
>>> oct(a)
'0o0443'
>>> b.oct
'0o7'
>>> (b + '0b0').oct
ValueError: Cannot convert to octal unambiguously - not multiple of 3 bits.
```

uint

To interpret the `BitString` as a binary (base-2) unsigned integer (i.e. a non-negative integer) use the `uint` property.

```
>>> a.uint
283
>>> b.uint
7
```

int

For a two's complement interpretation as a base-2 signed integer use the `int` property. If the first bit of the `BitString` is zero then the `int` and `uint` interpretations will be equal, otherwise the `int` will represent a negative number.

```
>>> a.int
283
>>> b.int
-1
```

data

A common need is to retrieve the raw bytes from a `BitString` for further processing or for writing to a file. For this use the `data` interpretation, which returns an ordinary Python string.

If the length of the `BitString` isn't a multiple of eight then it will be padded with between one and seven zero bits up to a byte boundary.

```
>>> open('somefile', 'wb').write(a.data)
>>> a2 = BitString(filename='somefile')
>>> a2.hex
'0x1230'
```

Note the extra four bits that were needed to byte align.

ue

The `ue` property interprets the `BitString` as a single unsigned exponential-Golomb code and returns an integer. If the `BitString` is not exactly one code then a `BitStringError` is raised instead. If you wish to read the next bits in the stream and interpret them as a code use the `readue` function. See Appendix B for a short explanation of this type of integer representation.

```
>>> s = BitString(ue=12)
>>> s.bin
'0b0001101'
>>> s.append(BitString(ue=3))
BitString('0x1a4')
>>> print s.readue(), s.readue()
12 3
```

se

The `se` property does much the same as `ue` and the provisos there all apply. The obvious difference is that it interprets the `BitString` as a signed exponential-Golomb rather than unsigned - see Appendix B in this tutorial for more information.


```
>>> s = BitString('0x164b')
>>> s.se
BitStringError: BitString is not a single exponential-Golomb code.
>>> while s.bitpos < s.length:
...     print s.readse()
-5
2
0
-1
```

3. Slicing, Dicing and Splicing

3.1. Slicing

Slicing can be done in couple of ways. The `slice` function takes three arguments: the first bit position you want, one past the last bit position you want and a multiplicative factor which defaults to 1. So for example `a.slice(10,12)` will return a 2-bit `BitString` of the 10th and 11th bits in `a`.

An equivalent method is to use indexing: `a[10:12]`. Note that as always the unit is bits, rather than bytes.

```
>>> a = BitString('0b00011110')
>>> b = a[3:7]
>>> c = a.slice(3, 7)           # s.slice(x, y) is equivalent to s[x:y]
>>> print a, b, c
0x1e 0xf 0xf
```

Indexing also works for missing and negative arguments, just as it does for other containers.

```
>>> a = BitString('0b00011110')
>>> print a[:5]                # first 5 bits
0b00011
>>> print a[3:]                # everything except first 3 bits
0b11110
>>> print a[-4:]               # final 4 bits
0xe
>>> print a[:-1]               # everything except last bit
0b0001111
>>> print a[-6:-4]             # from 6 from the end to 4 from the end
0b01
```

Stepping in slices

The step parameter (also known as the stride) can be used in slices. Its use is a little non-standard as it effectively gives a multiplicative factor to apply to the start and stop parameters, rather than skipping over bits.

For example this makes it much more convenient if you want to give slices in terms of bytes instead of bits. Instead of writing `s[a*8:b*8]` you can use `s[a:b:8]`.

When using a step, the `BitString` is effectively truncated to a multiple of the step, so `s[:8]` is equal to `s` if `s` is an integer number of bytes, otherwise it is truncated by up to 7 bits. So the final seven complete 16-bit words could be written as `s[-7::16]`.

```
>>> a = BitString('0x470000125e')
>>> print a[0:4:8]             # The first four bytes
0x47000012
>>> print a[-3::4]             # The final three nibbles
0x25e
```

Negative slices are also allowed, and should do what you'd expect. So for example `s[::-1]` returns a bit-reversed copy of `s` (which is similar to `s.reversebits()`, which does the same operation on `s` in-place). As another example, to get the first 10 bytes in reverse byte order you could use `s_bytereversed = s[0:10:-8]`.

```
>>> print a[:-5:-4]           # Final five nibbles reversed
0xe5210
>>> print a[::-8]             # The whole BitString byte reversed
0x5e12000047
```

3.2. Joining

To join together a couple of `BitString` objects use the `+` or `+=` operators, or the `append` and `prepend` functions.

```
# Six ways of creating the same BitString:

a1 = BitString(bin='000') + BitString(hex='f')
a2 = BitString('0b000') + BitString('0xf')
a3 = BitString('0b000') + '0xf'
a4 = BitString('0b000').append('0xf')
a5 = BitString('0xf').prepend('0b000')
a6 = BitString('0b000')
a6 += '0xf'
```

If you want to join a large number of `BitString` objects then the function `join` can be used to improve efficiency and readability.

```
# Don't do it this way!
s = BitString()
for bs in bslist:
    s = s + bs

# This is much more efficient:
s = bitstring.join(bslist)
```

3.3. Truncating

The truncate functions modify the `BitString` that they operate on, but also return themselves.

```
>>> a = BitString('0x001122')
>>> a.truncateend(8)
BitString('0x0011')
>>> b = a.truncatestart(8)
>>> a == b == '0x11'
True
```

A similar effect can be obtained using slicing - the major difference being that a new `BitString` is returned and the `BitString` being operated on remains unchanged.

3.4. Inserting, deleting and overwriting

`insert` takes one `BitString` and inserts it into another. A bit position can be specified, but if not present then the current `bitpos` is used.

```
>>> a = BitString('0x00112233')
>>> b = BitString('0xffff')
>>> a.insert(b, 16)
>>> a.hex
'0x0011ffff2233'
```

You can also use a string with `insert`, which will be interpreted as a binary or hexadecimal string. So the previous example could be written without using `b` as:

```
>>> a.insert('0xffff', 16)
```

`overwrite` does much the same as `insert`, but as you might expect the `BitString` object's data is overwritten by the new data.

```
>>> a = BitString('0x00112233')
>>> a.bitpos = 4
>>> a.overwrite('0b1111')           # Uses current bitpos as default
>>> a.hex
'0x0f112233'
```

`deletebits` and `deletebytes` remove sections of the `BitString`. By default they remove at the current `bitpos` - this must be at a byte boundary if using `deletebytes`:

```
>>> a = BitString('0b00011000')
>>> a.deletebits(2, 3)               # remove 2 bits at bitpos 3
>>> a.bin
'0b000000'

>>> b = BitString('0x112233445566')
>>> b.bytepos = 3
>>> b.deletebytes(2)
>>> b.hex
'0x11223366'
```

3.5. The BitString as a list

If you treat a `BitString` object as a list whose elements are all either '1' or '0' then you won't go far wrong. The table below gives some of the equivalent ways of using functions and the standard slice notation.

Using functions	Using slices
<code>s.truncatestart(bits)</code>	<code>del s[:bits]</code>
<code>s.truncateend(bits)</code>	<code>del s[-bits:]</code>
<code>s.slice(startbit, endbit, step)</code>	<code>s[startbit:endbit:step]</code>
<code>s.insert(bs, bitpos)</code>	<code>s[bitpos:bitpos] = bs</code>
<code>s.overwrite(bs, bitpos)</code>	<code>s[bitpos:bitpos+len(bs)] = bs</code>
<code>s.deletebits(bits, bitpos)</code>	<code>del s[bitpos:bitpos+bits]</code>
<code>s.deletebytes(bytes, bytepos)</code>	<code>del s[bytepos:bytepos+bytes:8]</code>
<code>s.append(bs)</code>	<code>s[len(s):len(s)] = bs</code>
<code>s.prepend(bs)</code>	<code>s[0:0] = bs</code>

3.6. Splitting

Sometimes it can be very useful to use a delimiter to split a `BitString` into sections. The `split` function returns a generator for the sections.

```
>>> a = BitString('0x4700004711472222')
>>> for s in a.split('0x47'):
...     print "Empty" if s.empty() else s.hex
Empty
0x470000
0x4711
0x472222
```

Note that the first item returned is always the `BitString` before the first occurrence of the delimiter, even if it is empty.

If you just want to split into equal parts then use the `cut` function. This takes a number of bits as its first argument and returns a generator for chunks of that size.

```
>>> a = BitString('0x47001243')
>>> for byte in a.cut(8):
...     print byte.hex
0x47
0x00
0x12
0x43
```

4. Reading and Navigating

4.1. Reading

A common need is to parse a large `BitString` into smaller syntax elements. Functions for reading in the `BitString` as if it were a file or stream are provided and will return new `BitString` objects. These new objects are top-level `BitString` objects and can be interpreted using properties as in the next example or could be read from themselves to form a hierarchy of reads.

Every `BitString` has a property `bitpos` which is the current position from which reads occur. `bitpos` can range from zero (its value on construction) to the length of the `BitString`, a position from which all reads will fail as it is past the last bit.

This example does some simple parsing of the supplied MPEG-1 video stream.

```
s = BitString(filename='test/test.mlv')
start_code = s.readbytes(4).hex
width = s.readbits(12).uint
height = s.readbits(12).uint
s.advancebits(37)
flags = s.readbits(2)
constrained_parameters_flag = flags.readbit().uint
load_intra_quantiser_matrix = flags.readbit().uint
```

In addition to the `read` functions there are matching `peek` functions. These are identical to the `read` except that they do not advance the position in the `BitString`.

```
s = BitString('0x4732aa34')
if s.peekbyte() == '0x47':
    t = s.readbytes(2)          # t.hex == '0x4732'
else:
    s.find('0x47')
```

The full list of functions is `readbit()`, `readbits(n)`, `readbyte()`, `readbytes(n)`, `peekbit()`, `peekbits(n)`, `peekbyte()` and `peekbytes(n)`.

4.2. Seeking

The properties `bitpos` and `bytepos` are available for getting and setting the position, which is zero on creation of the `BitString`. There are also `advance`, `retreat` and `seek` functions that perform equivalent actions. Whether you use the functions or the properties is purely a personal choice.

Note that you can only use `bytepos` or the `advance/retreatbyte(s)` functions if the position is byte aligned, i.e. the bit position is a multiple of 8. Otherwise a `BitStringError` is raised.

Using functions	Using properties
<code>advancebit()</code>	<code>bitpos += 1</code>
<code>advancebits(n)</code>	<code>bitpos += n</code>
<code>advancebyte()</code>	<code>bytepos += 1</code>
<code>advancebytes(n)</code>	<code>bytepos += n</code>
<code>retreatbit()</code>	<code>bitpos -= 1</code>
<code>retreatbits(n)</code>	<code>bitpos -= n</code>
<code>retreatbyte()</code>	<code>bytepos -= 1</code>
<code>retreatbytes(n)</code>	<code>bytepos -= n</code>
<code>seekbit(p)</code>	<code>bitpos = p</code>
<code>seekbyte(p)</code>	<code>bytepos = p</code>
<code>tellbit()</code>	<code>bitpos</code>
<code>tellbyte()</code>	<code>bytepos</code>

For example:

```
>>> s = BitString('0x123456')
>>> s.bitpos
0
>>> s.bytepos += 2
>>> s.bitpos           # note bitpos verses bytepos
16
>>> s.advancebits(4)
>>> print s.read(4).bin # the final nibble '0x6'
0b0110
```

4.3. Finding and replacing

To search for a sub-string use the `find` function. If the find succeeds it will set the position to the start of the next occurrence of the searched for string and return `True`, otherwise it will return `False`. By default the sub-string will only be found on byte boundaries; to allow it to be found at any position set `bytealigned=False`.

```
>>> s = BitString('0x00123400001234')
>>> found = s.find('0x1234')
>>> print found, s.bitpos
True 1
>>> found = s.find('0xff')
>>> print found, s.bitpos
False 1
```

`rfind` does much the same as `find`, except that it will find the last occurrence, rather than the first.

```
>>> t = BitString('0x0f231443e8')
>>> found = t.rfind('0xf', False)      # Search all bit positions in reverse
>>> print found, t.bitpos
True 31                                # Found within the 0x3e near the end
```

To find all occurrences of a `BitString` inside another (even overlapping ones), use `findall`. This returns a generator for the bit positions of the found strings.

```
>>> r = BitString('0b011101011001')
>>> ones = r.findall('0b1', bytealigned=False)
>>> print list(ones)
[1, 2, 3, 5, 7, 8, 11]
```

To replace all occurrences of one `BitString` with another use `replace`.

```
>>> s = BitString('0b110000110110')
>>> s.replace('0b110', '0b1111', bytealigned=False)
3                                     # The number of replacements made
>>> s.bin
'0b111100011111111'
```

5. Miscellany

5.1. Other Functions

empty()

Returns `True` if the `BitString` contains no data (i.e. has zero length). Otherwise returns `False`.

```
>>> a = BitString()
>>> print a.empty()
True
```

bytealign()

This function advances between zero and seven bits to make the `bitpos` a multiple of eight. It returns the number of bits advanced.

```
>>> a = BitString('0x11223344')
>>> a.bitpos = 1
>>> skipped = a.bytealign()
>>> print skipped, a.bitpos
7 8
>>> skipped = a.bytealign()
>>> print skipped, a.bitpos
0 8
```

reversebits()

This simply reverses the bits of the `BitString` in place and returns `self`. You can optionally specify a range of bits to reverse.

```
>>> a = BitString('0b000001101')
>>> a.reversebits().bin
'0b101100000'
>>> a.reversebits(0, 4).bin
'0b110100000'
```

5.2. Special Methods

A few of the special methods have already been covered, for example `__add__` and `__iadd__` (the `+` and `+=` operators) and `__getitem__` and `__setitem__` (reading and setting slices via `[]`). Here are the rest:

`__len__`

This implements the `len` function and returns the length of the `BitString` in bits. There's not much more to say really, except to emphasise that it is always in bits and never bytes.

```
>>> len(BitString('0x00'))
8
```

`__str__` , `__repr__`

These get called when you try to print a `BitString`. As `BitString` objects have no preferred interpretation the form printed might not be what you want - if not then use the `hex`, `bin`, `int` etc. properties. The main use here is in interactive sessions when you just want a quick look at the `BitString`. The `__repr__` tries to give a code fragment which if evaluated would give an equal `BitString`.

The form used for the `BitString` is generally the one which gives it the shortest representation. If the resulting string is too long then it will be truncated with `'...'` - this prevents very long `BitString` objects from tying up your interactive session printing themselves.

```
>>> a = BitString('0b1111 111')
>>> print a
0b1111111
>>> a
BitString('0b1111111')
>>> a += '0b1'
>>> print a
0xff
>>> print a.bin
0b11111111
```

`__eq__` , `__ne__`

The equality of two `BitString` objects is determined by their binary representations being equal. If you have a different criterion you wish to use then code it explicitly, for example `a.int == b.int` could be true even if `a == b` wasn't (as they could be different lengths).

Note that two `BitString` objects can have different offsets, but still be equal if their binary representations are equal.

```
>>> BitString('0b0010') == '0x2'
True
>>> BitString('0x2') != '0o2'
True
```

`__hex__` , `__oct__`

You can if you wish use the built-in functions `hex()` and `oct()` instead of the `hex` and `oct` properties, although for consistency it is probably better to stick to using the properties. Note that octals are always prefixed by `'0o'` rather than just `'0'`. Note also that although a `bin()` built-in function was introduced in

Python 2.6 there doesn't seem to be a corresponding `__bin__` special function, for reasons that escape me, so you can't use `bin()` on a `BitString`.

```
>>> a = BitString('0o7777')
>>> a.oct
'0o7777'
>>> oct(a)
'0o7777'
>>> a.hex
'0xffff'
>>> hex(a)
'0xffff'
```

`__invert__`

To invert all the bits in a `BitString` use the `~` operator.

```
>>> a = BitString('0b0001100111')
>>> print a
0b0001100111
>>> print ~a
0b1110011000
>>> ~~a == a
True
```

`__lshift__` , `__rshift__` , `__ilshift__` , `__irshift__`

Bitwise shifts can be achieved using `<<`, `>>`, `<=<` and `>=>`. Bits shifted off the left or right are replaced with zero bits. If you need special behaviour, such as keeping the sign of two's complement integers then do the shift on the property instead.

```
>>> a = BitString('0b10011001')
>>> b = a << 2
>>> print b
0b01100100
>>> a >>= 2
>>> print a
0b00100110
```

`__mul__` , `__imul__` , `__rmul__`

Multiplication of a `BitString` by an integer means the same as it does for ordinary strings: concatenation of multiple copies of the `BitString`.

```
>>> a = BitString('0b10')*8
>>> print a.bin
0b1010101010101010
```

`__copy__`

This allows the `BitString` to be copied via the `copy` module.

```
>>> import copy
>>> a = BitString('0x4223fbddec2231')
>>> b = copy.copy(a)
>>> b == a
True
>>> b is a
False
```

It's not terribly exciting, and isn't even the preferred method of making a copy. Using `b = BitString(a)` is another option, but `b = a[:]` may be more familiar to some.

`__and__` , `__or__` , `__xor__`

Bit-wise AND, OR and XOR are provided for `BitString` objects of equal length only (otherwise a `ValueError` is raised). The right-hand-side of expression can be a string to use in the auto initialiser.

```
>>> a = BitString('0b00001111')
>>> b = BitString('0b01010101')
>>> print (a&b).bin
0b00000101
>>> print (a|b).bin
0b01011111
>>> print (a^b).bin
0b01010000
>>> b &= '0x1f'
>>> print b.bin
0b00010101
```

A. Reference

The `bitstring` module provides just one class, `BitString`, whose public methods, special methods and properties are detailed in this section.

Note that in places where a `BitString` can be used as a parameter, any other valid input to the auto initialiser can also be used. This means that the parameter can be any one of:

- A `BitString`.
- A string starting with `0x`, interpreted as hexadecimal.
- A string starting with `0o`, interpreted as octal.
- A string starting with `0b`, interpreted as binary.
- A list or tuple, whose elements will be evaluated as booleans, and the bits set to 1 for `True` items and 0 for `False` items.

A.1. BitString methods and properties

advancebit

`s.advancebit()`

Advances position by 1 bit. Equivalent to `s.bitpos += 1`.

advancebits

`s.advancebits(bits)`

Advances position by `bits` bits.

Equivalent to `s.bitpos += bits`.

advancebyte

`s.advancebyte()`

Advances position by 8 bits.

Equivalent to `s.bitpos += 8`.

Unlike the alternative, `s.bytepos += 1`, `advancebyte` will not raise a `BitStringError` if the current position is not byte-aligned.

advancebytes

`s.advancebytes(bytes)`

Advances position by `8*bytes` bits.

Equivalent to `s.bitpos += 8*bytes`. Unlike the alternative, `s.bytepos += bytes`, `advancebytes` will not raise a `BitStringError` if the current position is not byte-aligned.

append

`s.append(bs)`

Join a `BitString` to the end of the current `BitString`. Returns `self`.

```
>>> s = BitString('0xbad')
>>> s.append('0xf00d')
BitString('0xbadf00d')
```

bin

`s.bin`

Read and write property for setting and getting the representation of the `BitString` as a binary string starting with `'0b'`.

When used as a getter, the returned value is always calculated - the value is never cached. When used as a setter the length of the `BitString` will be adjusted to fit its new contents.

```
if s.bin == '0b001':
    s.bin = '0b1111'

# Equivalent to s.append('0b1')
s.bin += '1'
```

bitpos

`s.bitpos`

Read and write property for setting and getting the current bit position in the `BitString`. Can be set to any value from 0 to length.

```
if s.bitpos < 100:
    s.bitpos += 10
```

bytealign

`s.bytealign()`

Aligns to the start of the next byte (so that `bitpos` is a multiple of 8) and returns the number of bits skipped.

If the current position is already byte aligned then it is unchanged.

```
>>> s = BitString('0xabcdef')
>>> s.advancebits(3)
>>> s.bytealign()
5
>>> s.bitpos
8
```

bytepos

s.bytepos

Read and write property for setting and getting the current byte position in the `BitString`.

When used as a getter will raise a `BitStringError` if the current position is not byte aligned.

cut

s.cut(bits, startbit=None, endbit=None, count=None)

Returns a generator for slices of the `BitString` of length bits.

At most count items are returned and the range is given by the slice `[startbit:endbit]`, which defaults to the whole `BitString`.

```
>>> s = BitString('0x1234')
>>> for nibble in s.cut(4):
...     s.prepend(nibble)
>>> print s
0x43211234
```

data

s.data

Read and write property for setting and getting the underlying byte data that contains the `BitString`.

Set using an ordinary Python string - the length will be adjusted to contain the data. When used as a getter the `BitString` will be padded with between zero and seven '0' bits to make it byte aligned.

The data property can be used to output your `BitString` to a file - just open a file in binary write mode and write the `BitString`'s data.

```
>>> s = BitString()
>>> s.data = 'hello'
>>> s.hex
'0x68656c6c6f'
>>> s += '0b01'
>>> s.data
'hello@'
```

deletebits

s.deletebits(bits, bitpos=None)

Removes bits from the `BitString` at position bitpos and returns self. If bitpos is not specified then the current position is used.

```
>>> s = BitString('0b1111001')
>>> s.deletebits(2, 4)
BitString('0b1111')
```

deletebytes

s.deletebytes(bytes, bytepos=None)

Removes bytes from the `BitString` at position bytepos and returns self.

If bytepos is not specified then the current position is used, provided it is byte aligned, otherwise `BitStringError` is raised.

empty

s.empty()

Returns True if the `BitString` is empty, i.e. has length==0. Otherwise returns False.

find

s.find(bs, bytealigned=True, startbit=None, endbit=None)

Searches for bs in the current `BitString` and sets bitpos to the start of bs and returns True if found, otherwise it returns False.

If bytealigned is True then it will look for bs only at byte aligned positions (which is generally much faster than searching for it in every possible bit position). startbit and endbit give the search range and default to 0 and s.length respectively.

findall

s.findall(bs, bytealigned=True, startbit=None, endbit=None, count=None)

Searches for all occurrences of bs (even overlapping ones) and returns a generator of their bit positions.

If bytealigned is True then bs will only be looked for at byte aligned positions. startbit and endbit optionally define a search range and default to 0 and s.length respectively.

hex

s.hex

Read and write property for setting and getting the hexadecimal representation of the `BitString`.

When used as a getter the value will be preceded by 0x, which is optional when setting the value. If the `BitString` is not a multiple of four bits long then getting its hex value will raise a `ValueError`.

```
>>> s = BitString(bin='1111 0000')
>>> s.hex
'0xf0'
>>> s.hex = 'abcdef'
>>> s.hex
'0xabcdef'
```

insert

s.insert(bs, bitpos=None)

Inserts bs at bitpos and returns self. The default for bitpos is the current position.

int

s.int

Read and write property for setting and getting the signed two's complement integer representation of the BitString. When used as a setter the value must fit into the current length of the BitString, else a ValueError will be raised.

```
>>> s = BitString('0xf3')
>>> s.int
-13
>>> s.int = 1232
ValueError: int 1232 is too large for a
BitString of length 8.
```

length

s.length

Read-only property that gives the length of the BitString in bits.

oct

s.oct

Read and write property for setting and getting the octal representation of the BitString. When used as a getter the value will be preceded by 0o, which is optional when setting the value. If the BitString is not a multiple of three bits long then getting its oct value will raise a ValueError.

```
>>> s = BitString('0b111101101')
>>> s.oct
'0o755'
>>> s.oct = '01234567'
>>> s.oct
'0o01234567'
```

overwrite

s.overwrite(bs, bitpos=None)

Replaces the contents of the current BitString with bs at bitpos and returns self. The default for bitpos is the current position.

```
>>> s = BitString(length=10)
>>> s.overwrite('0b111', 3)
BitString('0b0001110000')
```

peekbit

s.peekbit()

Returns the next bit in the current BitString as a new BitString but does not advance the position.

peekbits

s.peekbits(bits)

Returns the next bits bits of the current BitString as a new BitString but does not advance the position.

```
>>> s = BitString('0xf01')
>>> s.bitpos = 4
>>> s.peekbits(4)
BitString('0x0')
>>> s.peekbits(8)
BitString('0x01')
```

peekbyte

s.peekbyte()

Returns the next byte of the current BitString as a new BitString but does not advance the position.

peekbytes

s.peekbytes(bytes)

Returns the next bytes bytes of the current BitString as a new BitString but does not advance the position.

prepend

s.prepend(bs)

Inserts bs at the beginning of the current BitString. Returns self.

```
>>> s = BitString('0b0').prepend('0xf')
>>> s
BitString('0b11110')
```

readbit

s.readbit()

Returns the next bit of the current BitString as a new BitString and advances the position.

readbits

s.readbits(bits)

Returns the next *bits* bits of the current BitString as a new BitString and advances the position.

```
>>> s = BitString('0x0001e2')
>>> s.readbits(16)
BitString('0x0001')
>>> s.readbits(3).bin
'0b111'
>>> s.bitpos
19
```

readbyte

s.readbyte()

Returns the next byte of the current BitString as a new BitString and advances the position.

readbytes

s.readbytes(bytes)

Returns the next *bytes* bytes of the current BitString as a new BitString and advances the position.

readse

s.readse()

Reads an exponential Golomb code from the stream at the current *bitpos* and returns its signed integer interpretation.

```
>>> s = BitString(se=-9) + BitString(se=4)
>>> print s
0x0988
>>> s.readse()
-9
>>> s.readse()
4
```

readue

s.readue()

Reads an exponential Golomb code from the stream at the current *bitpos* and returns its unsigned integer interpretation.

replace

s.replace(old, new, bytealigned=True, startbit=None, endbit=None, count=None)

Finds occurrences of *old* and replaces them with *new*. Returns the number of replacements made.

If *bytealigned* is *True* then replacements will only be made on byte boundaries. *startbit* and *endbit* give the search range and default to 0 and *s.length* respectively. If *count* is specified then no more than this many replacements will be made.

retreatbit

s.retreatbit()

Retreats position by 1 bit. Equivalent to *bitpos -= 1*.

retreatbits

s.retreatbits(bits)

Retreats position by *bits* bits. Equivalent to *bitpos -= bits*.

retreatbyte

s.retreatbyte()

Retreats position by 8 bits. Equivalent to *bitpos -= 8*. Unlike the alternative, *bytepos -= 1*, *retreatbyte* will not raise a *BitStringError* if the current position is not byte-aligned.

retreatbytes

s.retreatbytes(bytes)

Retreats position by *bytes**8 bits. Equivalent to *bitpos -= 8*bytes*. Unlike the alternative, *bytepos -= bytes*, *retreatbytes* will not raise a *BitStringError* if the current position is not byte-aligned.

reversebits

s.reversebits(startbit=None, endbit=None)

Reverses bits in the BitString in-place and returns self. *startbit* and *endbit* give the range and default to 0 and *s.length* respectively.

```
>>> a = BitString('0b10111')
>>> a.reversebits().bin
'0b11101'
```

rfind

```
s.rfind(bs, bytealigned=True,
        startbit=None, endbit=None)
```

Searches backwards for `bs` in the current `BitString` and returns `True` if found.

If `bytealigned` is `True` then it will look for `bs` only at byte aligned positions. `startbit` and `endbit` give the search range and default to 0 and `self.length` respectively. Note that as it's a reverse search it will start at `endbit` and finish at `startbit`.

se

```
s.se
```

Read and write property for setting and getting the signed exponential-Golomb code representation of the `BitString`. The property is set from an signed integer, and when used as a getter a `BitStringError` will be raised if the `BitString` is not a single code.

```
>>> s = BitString(se=-40)
>>> s.bin
0b0000001010001
>>> s += '0b1'
>>> s.se
BitStringError: BitString is not a single
exponential-Golomb code.
```

seekbit

```
s.seekbit(bitpos)
```

Moves the current position to `bitpos`. Equivalent to `s.bitpos = bitpos`.

seekbyte

```
s.seekbyte(bytepos)
```

Moves the current position to `bytepos`. Equivalent to `s.bytepos = bytepos`.

slice

```
s.slice(startbit, endbit, step)
```

Returns the `BitString` slice `s[startbit*step : endbit*step]`.

The `step` parameter gives a multiplicative factor for the start and end positions, so for example using a `step` of 8 allows the slice to be given in terms of byte indices rather than bit indices.

It's use is equivalent to using the slice notation; see `__getitem__` for examples.

split

```
s.split(delimiter, bytealigned=True,
        startbit=None, endbit=None,
        count=None)
```

Splits `s` into sections that start with `delimiter`. Returns a generator for `BitString` objects.

The first item generated is always the bits before the first occurrence of `delimiter` (even if empty). A slice can be optionally specified with `startbit` and `endbit`, while `count` specifies the maximum number of items generated.

tellbit

```
s.tellbit()
```

Returns the current bit position.

Equivalent to using the `bitpos` property as a getter.

tellbyte

```
s.tellbyte()
```

Returns the current byte position.

Equivalent to using the `bytepos` property as a getter, and will raise a `BitStringError` if the `BitString` is not byte aligned.

truncateend

```
s.truncateend(bits)
```

Remove the last `bits` bits from the end of the `BitString`. Returns `self`.

truncatestart

```
s.truncatestart(bits)
```

Remove the first `bits` bits from the start of the `BitString`. Returns `self`.

ue

```
s.ue
```

Read and write property for setting and getting the unsigned exponential-Golomb code representation of the `BitString`.

The property is set from an unsigned integer, and when used as a getter a `BitStringError` will be raised if the `BitString` is not a single code.

uint

s.uint

Read and write property for setting and getting the unsigned base-2 integer representation of the BitString.

When used as a setter the value must fit into the current length of the BitString, else a ValueError will be raised.

A.2. BitString special methods

__add__ / __radd__

s1 + s2

Concatenate two BitString and return the result. Either s1 or s2 can be auto initialised.

```
s = BitString(ue=132) + '0xff'
s2 = '0b101' + s
```

__and__ / __rand__

s1 & s2

Returns the bit-wise AND between s1 and s2, which must have the same length otherwise a ValueError is raised.

```
>>> print BitString('0x33') & '0x0f'
0x03
```

__contains__

bs in s

Returns True if bs can be found in s, otherwise returns False.

Equivalent to using find with bytealigned=False, except that bitpos will not be changed.

```
>>> '0b11' in BitString('0x06')
True
>>> '0b111' in BitString('0x06')
False
```

__copy__

s2 = copy.copy(s1)

This allows the copy module to correctly copy BitString objects. Other equivalent methods are to initialise a new BitString with the old one or to take a complete slice.

```
>>> import copy
>>> s = BitString('0o775')
>>> s_copy1 = copy.copy(s)
>>> s_copy2 = BitString(s)
>>> s_copy3 = s[:]
>>> s == s_copy1 == s_copy2 == s_copy3
True
```

__delitem__

del s[startbit:endbit:step]

Deletes the slice specified.

__eq__

s1 == s2

Compares two BitString objects for equality, returning True if they have the same binary representation, otherwise returning False.

```
>>> BitString('0o7777') == '0xffff'
True
>>> a = BitString(uint=13, length=8)
>>> b = BitString(uint=13, length=10)
>>> a == b
False
```

__getitem__

s[startbit:endbit:step]

Returns a slice of s.

The usual slice behaviour applies except that the step parameter gives a multiplicative factor for startbit and endbit (i.e. the bits 'stepped over' are included in the slice).

```
>>> s = BitString('0x0123456')
>>> s[0:4]
BitString('0x1')
>>> s[0:3:8]
BitString('0x012345')
```

__hex__

hex(s)

Returns the hexadecimal representation of the BitString, i.e. a string starting with 0x.

Equivalent to using the hex property, and so will raise a ValueError if the BitString is not a multiple of four bits long.

__iadd__

s1 += s2

Append a BitString to the current BitString and return the result.

```
>>> s = BitString(ue=423)
>>> s += BitString(ue=12)
>>> s.readue()
423
>>> s.readue()
12
```

__ilshift__

s <<= n

Shifts the bits in *s* in place to the left by *n* places. Returns self. Bits shifted off the left hand side are lost, and replaced by 0 bits on the right hand side.

__imul__

*s *= n*

Concatenates *n* copies of *s* and returns self. Raises ValueError if *n* < 0.

```
>>> s = BitString('0xef')
>>> s *= 3
>>> print s
0xefefef
```

__init__

s = BitString(auto=None, length=None, offset=0, data=None, filename=None, hex=None, bin=None, oct=None, uint=None, ue=None, se=None)

Creates a new BitString. You must specify at most one of the initialisers *auto*, *data*, *bin*, *hex*, *oct*, *uint*, *int*, *se*, *ue* or *filename*. If no initialiser is given then a zeroed BitString of length *bits* is created.

offset is optional for most initialisers, but only really useful for *data* and *filename*. It gives a number of bits to ignore at the start of the BitString.

Specifying *length* is mandatory when using the *int* and *uint* initialisers. It must be large enough that a BitString can contain the integer in *length* bits. It is an error to specify *length* when using the *ue* or *se* initialisers. For other initialisers *length* can be used to truncate data from the end of the input value.

```
>>> s1 = BitString(hex='0x934')
>>> s2 = BitString(oct='0o4464')
>>> s3 = BitString(bin='0b001000110100')
>>> s4 = BitString(int=-1740, length=12)
>>> s5 = BitString(uint=2356, length=12)
>>> s6 = BitString(data='\x93@', length=12)
>>> s1 == s2 == s3 == s4 == s5 == s6
True
```

__invert__

~s

Returns the BitString with every bit inverted, that is all zeros replaced with ones, and all ones replaced with zeros. If the BitString is empty then a BitStringError will be raised.

```
>>> s = BitString('0b1110010')
>>> print ~s
0b0001101
>>> print ~s & s
0b0000000
```

__irshift__

s >>= n

Shifts the bits in *s* in place by *n* places to the right and returns self. The *n* left-most bits will become zeros.

```
>>> s = BitString('0b110')
>>> s >>= 2
>>> s.bin
'0b001'
```

__len__

len(s)

Returns the length of the BitString in bits.

__lshift__

s << n

Returns the BitString with its bits shifted *n* places to the left (or *len(s)* if it's less). The *n* right-most bits will become zeros.

```
>>> s = BitString('0xff')
>>> s << 4
BitString('0xf0')
```

__mul__ / **__rmul__**

*s * n* / *n * s*

Return BitString consisting of *n* concatenations of *s*.


```
>>> a = BitString('0x34')
>>> b = a*5
>>> print b
0x3434343434
```

__ne__

```
s1 != s2
```

Compares two `BitString` objects for inequality, returning `False` if they have the same binary representation, otherwise returning `True`.

__oct__

```
oct(s)
```

Returns the octal representation of the `BitString`, i.e. a string starting with `0o`.

Equivalent to using the `oct` property, and so will raise a `ValueError` if the `BitString` is not a multiple of three bits long.

__or__ / **__ror__**

```
s1 | s2
```

Returns the bit-wise OR between `s1` and `s2`, which must have the same length otherwise a `ValueError` is raised.

```
>>> print BitString('0x33') | '0x0f'
0x3f
```

__repr__

```
repr(s)
```

A representation of the `BitString` that could be used to create it (which will often not be the form used to create it).

If the result is too long then it will be truncated with `'...'` and the length of the whole `BitString` will be given.

```
>>> BitString('0b11100011')
BitString('0xe3')
```

__rshift__

```
s >> n
```

Returns the `BitString` with its bits shifted `n` places to the right (or `len(s)` if it's less). The `n` left-most bits will become zeros.

```
>>> s = BitString('0xff')
>>> s >> 4
BitString('0x0f')
```

__setitem__

```
s1[startbit:endbit:step] = s2
```

Replaces the slice specified with `s2`.

```
>>> s = BitString('0x00112233')
>>> s[1:2:8] = '0xfff'
>>> print s
0x00fff2233
>>> s[-12:] = '0xc'
>>> print s
0x00fff2c
```

__str__

```
print s
```

Prints a representation of `s`, trying to be as brief as possible.

If `s` is a multiple of 4 bits long then hex will be used, otherwise if it's a multiple of 3 bits then octal is used, otherwise binary. Very long strings will be truncated with `'...'`.

```
>>> s = BitString('0b1')*6
>>> print s
0o77
>>> print s + '0b1'
0b1111111
>>> print s + '0b11'
0xff
```

__xor__ / **__rxor__**

```
s1 ^ s2
```

Returns the bit-wise XOR between `s1` and `s2`, which must have the same length otherwise a `ValueError` is raised. Either `s1` or `s2` can be a string for the auto initialiser.

```
>>> print BitString('0x33') ^ '0x0f'
0x3c
```

A.3. Module functions

join

```
s = bitstring.join(bsl)
```

Returns the concatenation of the `BitString` objects in the list `bsl`.

```
>>> s = join(['0x0001ee', BitString(int=13,
length=24), '0b0111'])
>>> print s
0x0001ee00000d7
```

B. Exponential-Golomb Codes

As this type of representation of integers isn't as well known as the standard base-2 representation I thought that a short explanation of them might be welcome. This section can be safely skipped if you're not interested.

Exponential-Golomb codes represent integers using bit patterns that get longer for larger numbers. For unsigned and signed numbers (the `BitString` properties `ue` and `se` respectively) the patterns start like this:

Bit pattern	Unsigned	Signed
1	0	0
010	1	1
011	2	-1
00100	3	2
00101	4	-2
00110	5	3
00111	6	-3
0001000	7	4
0001001	8	-4
0001010	9	5
0001011	10	-5
0001100	11	6
...

They consist of a sequence of n '0' bits, followed by a '1' bit, followed by n more bits. The bits after the first '1' bit count upwards as ordinary base-2 binary numbers until they run out of space and an extra '0' bit needs to get included at the start.

The advantage of this method of representing integers over many other methods is that it can be quite efficient at representing small numbers without imposing a limit on the maximum number that can be represented.

Exercise: Using the table above decode this sequence of unsigned Exponential Golomb codes:

```
001001101101101011000100100101
```

The answer is that it decodes to 3, 0, 0, 2, 2, 1, 0, 0, 8, 4. Note how you don't need to know how many bits are used for each code in advance - there's only one way to decode it. To create this bitstring you could have written something like:

```
a = bitstring.join([BitString(ue=i) for i in [3,0,0,2,2,1,0,0,8,4]])
```

and to read it back:

```
while a.bitpos != a.length:  
    print a.readue()
```

The notation `ue` and `se` for the exponential-Golomb code properties comes from the H.264 video standard, which uses these types of code a lot. The particular way that the signed integers are represented might be peculiar to this standard as I haven't seen it elsewhere (and an obvious alternative is minus the one given here), but the unsigned mapping seems to be universal.

C. Internals

I am including some information on the internals of the `BitString` class here, things that the general user shouldn't need to know. The objects and methods described here all start with an underscore, which means that they are a private part of the implementation, not a part of the public interface and that that I reserve the right to change, rename and remove them at any time!

The data in a `BitString` can be considered to consist of three parts.

- The byte data, either contained in memory, or as part of a file.
- A length in bits.
- An offset to the data in bits.

Storing the data in byte form is pretty essential, as anything else could be very memory inefficient. Keeping an offset to the data allows lots of optimisations to be made as it means that the byte data doesn't need to be altered for almost all operations. An example is in order:

```
a = BitString('0x01ff00')
b = a[7:12]
```

This is about as simple as it gets, but let's look at it in detail. First `a` is created by parsing the string as hexadecimal (as it starts with `0x`) and converting it to three data bytes `\x01\xff\x00`. By default the length is the bit length of the whole string, so it's 24 in this case, and the offset is zero.

Next, `b` is created from a slice of `a`. This slice doesn't begin or end on a byte boundary, so one way of obtaining it would be to copy the data in `a` and start doing bit-wise shifts to get it all in the right place. This can get really very computationally expensive, so instead we utilise the offset and length parameters.

The procedure is simply to copy the byte data containing the substring and set the offset and length to get the desired result. So in this example we have:

```
a : data = '\x01\xff\x00', offset = 0, length = 24
b : data = '\x01\xff', offset = 7, length = 5
```

This method also means that `BitString` objects initialised from a file don't have to copy anything into memory - the data instead is obtained with a byte offset into the file. This brings us onto the different types of datastores used.

The `BitString` has a `_datastore` member, which at present is either a `_MemArray` class or a `_FileArray` class. The `_MemArray` class is really just a light wrapper around an `array.array` object that contains the real byte data, so when we were talking about the data earlier I was really referring to the byte data contained in the `array.array`, in the `_MemArray`, in the `_datastore`, in the `BitString` (but that seemed a bit much to give you in one go).